



Measuring Chain of Thought Faithfulness

Based on *Measuring Chain of Thought Faithfulness by Unlearning Reasoning Steps*

Group27

Yanpeng Liang Yonggong Teng

Dechao Huang Yao Yu

2026-5-28



復旦大學
FUDAN UNIVERSITY



We implemented the work related to “Measuring Chain of Thought Faithfulness”.
It is mainly based on the reproduction and extension of the paper *Measuring Chain of Thought Faithfulness by Unlearning Reasoning Steps*.



Outline

1 Introduction

- ▶ Introduction
- ▶ Reproduction and Extension Experiments
 - Reproduction
 - Chinese Reproduction
 - Extension Experiments1
 - Extension Experiments2
- ▶ Further Exploration
 - Measuring Tree of Thought Faithfulness
 - Distilled CoT via deepseek-v4-pro
- ▶ Conclusion, Limitations and Discussion

What is Chain of Thought Faithfulness

1 Introduction



Chain-of-Thought Prompting Elicits Reasoning in Large Language Models[1]



Measuring Faithfulness in Chain-of-Thought Reasoning[2]



Measuring Chain of Thought Faithfulness by Unlearning Reasoning Steps[3]





Parametric Faithfulness Framework (PFF)

1 Introduction

Parametric Faithfulness Framework (PFF) is a two-stage framework for measuring whether generated reasoning faithfully reflects the model's internal, parameter-level reasoning process.

Stage 1: Parameter Intervention

- Prompt the model to generate a Chain-of-Thought (CoT).
- Segment the CoT into individual reasoning steps.
- Apply a parameter-level intervention to remove information contained in each reasoning step.
- Verify that the intervention is valid: **Efficacy, Specificity, General Capabilities**

Stage 2: Faithfulness Evaluation

- Compare the predictions of the original model and the intervened model.
- If removing a reasoning step changes the model's prediction, the step is considered parametrically faithful.

5/52 ■ Optionally compare the generated reasoning before and after .



FUR: An Instantiation of PFF

1 Introduction

Faithfulness by Unlearning Reasoning Steps (FUR) instantiates PFF by using machine unlearning as the parameter intervention.

What FUR Does

- Generate a Chain-of-Thought (CoT) for each question.
- Segment the CoT into individual reasoning steps.
- Unlearn each reasoning step separately from the model parameters.
- Check whether the model's direct answer changes after unlearning.

Main Contribution

- FUR moves from **context-level perturbation** to **parameter-level intervention**.
- It uses **NPO + KL regularization** to erase information in reasoning steps.
- It quantifies faithfulness with:
 - **FF-HARD**: whether any step changes the final answer.
 - **FF-SOFT**: how much each step reduces the original answer probability.



NPO + KL: Machine Unlearning for Reasoning Steps

1 Introduction

NPO + KL is used in FUR as a **machine unlearning method** to erase information contained in a reasoning step from the model parameters.

Forget and Retain Sets

- **Forget set** $\mathcal{D}_{FG}$: content words from the target reasoning step.
- **Retain set** $\mathcal{D}_{RT}$: content words from unrelated CoT steps.

Negative Preference Optimization (NPO)

$$\mathcal{L}_{\text{NPO}} = -\frac{2}{\beta} \mathbb{E}_{(x,y) \sim \mathcal{D}_{FG}} \left[\log \sigma \left(-\beta \log \frac{\pi_{\theta}(y|x)}{\pi_{\text{ref}}(y|x)} \right) \right]$$

Final Objective

$$\mathcal{L} = \mathcal{L}_{\text{NPO}} + \lambda \mathcal{L}_{\text{KL Regularization}}$$



Outline

2 Reproduction and Extension Experiments

- ▶ Introduction
- ▶ Reproduction and Extension Experiments
 - Reproduction
 - Chinese Reproduction
 - Extension Experiments1
 - Extension Experiments2
- ▶ Further Exploration
 - Measuring Tree of Thought Faithfulness
 - Distilled CoT via deepseek-v4-pro
- ▶ Conclusion, Limitations and Discussion



Experimental Setup: Models and Datasets

2 Reproduction and Extension Experiments

Evaluation Models

4 representative instruction-tuned models across 3 families:

- **Meta LLaMA Family**
 - LLaMA-3-8B-Instruct
 - LLaMA-3.2-3B-Instruct
- **Mistral AI**
 - Mistral-7B-Instruct-v0.2
- **Microsoft**
 - Phi-3-mini-4k-Instruct

Unified Protocol: All models are evaluated in a **zero-shot** setting

Benchmark Datasets

5 diverse multi-hop multi-choice QA (MCQA) datasets:

- **ARC-Challenge:** Scientific reasoning
- **OpenbookQA:** Elementary science commonsense
- **StrategyQA:** Implicit multi-hop logic
- **TruthfulQA:** Factuality and common misconceptions
- **BigBench-Hard (Sports):** Sports domain multiple choice




Experimental Setup

2 Reproduction and Extension Experiments

Item	Setting
Models	LLaMA-3.2-3B-Instruct, LLaMA-3-8B-Instruct, Mistral-7B-Instruct-v0.2
Datasets	ARC-Challenge, OpenBookQA, StrategyQA
Task format	Multi-choice question answering (MCQA)
Prompting	Zero-shot, two-step CoT generation and answer extraction
Unlearning method	NPO + KL, stepwise unlearning
CoT cache	250 CoTs
instances; Specificity split	230 instances; 20 held-out samples
Updated parameters	FF2 / MLP down projection
Evaluation	faithfulness(ff_hard, main)
Metrics	Efficacy, Specificity, FF-HARD, FF-SOFT

Paper-LR Runs: Source-Level Reproduction


2 Reproduction and Extension Experiments



Main observation: when our learning rate matches the source-paper setting, hard faithfulness is mostly reproducible.

LR-Mismatched Runs: Older Results Move Differently

2 Reproduction and Extension Experiments




Main observation: The LR-mismatched comparisons suggest that the measured faithfulness signal is highly sensitive to optimization hyperparameters. Lower-than-paper learning rates tend to produce conservative under-unlearning. In contrast, more aggressive settings can inflate FF-hard by forcing large probability shifts, but this often comes with reduced specificity and less stable behavior.




FF-hard Gap and FF-soft Movement

2 Reproduction and Extension Experiments

Hard faithfulness gap




Soft probability movement



- Paper-LR runs usually have small FF-hard gaps; LR-mismatched runs produce the largest deviations.
- FF-soft separates two failure modes: weak probability movement for low-LR Phi-3, versus aggressive probability movement and specificity loss for high-LR StrategyQA runs.

FF-soft Example: One Step Across Five Epochs

2 Reproduction and Extension Experiments



- Case-level alignment changes with the learning-rate regime. Exact paper-LR comparisons show more shared flips between our run and the source baseline, while LR-mismatched settings produce more asymmetric flips. This suggests that hyperparameters affect not only the aggregate FF-hard score, but also which individual CoT steps appear causally important.



Subsection 2

Chinese Reproduction



Chinese Reproduction: Qwen3-8B + C-Eval

2 Reproduction and Extension Experiments

Experimental setup

Model	Qwen3-8B
Dataset	Chinese C-Eval sub-set
CoT cache	250 CoTs
Specificity split	20 held-out samples
Unlearning	NPO-KL, stepwise, sentencize
Updated parameters	FF2 / MLP down projection
Result scale	230 instances, 498 steps

Main quantitative results

Metric	Value
Faithful instances	21 / 230 = 9.13%
Per-step answer flips	25 / 498 = 5.02%
Mean efficacy	71.50%
Mean specificity	96.84%
Mean answer mass shift	0.0336

Interpretation. The unlearning intervention is effective and local, but most visible CoT steps do not causally determine the final answer. Strong faithful steps exist, but they are sparse.


Strong Positive Cases: Answer Probability Transfer

2 Reproduction and Extension Experiments

Top answer-flip examples

Instance	Step	Change	Shift
basic-med-10	1	C $\rightarrow$ B	1.0000
sports-sci-1	2	C $\rightarrow$ D	1.0000
plant-prot-10	0	C $\rightarrow$ B	0.99999
ms-history-18	0	A $\rightarrow$ B	0.99994
chinese-lang-19	4	D $\rightarrow$ B	0.99945


In these cases, unlearning one CoT step moves probability mass almost completely from the original option to another option, while specificity remains high. These are the clearest faithful reasoning-step cases.




Global Distribution and Step Saliency

2 Reproduction and Extension Experiments

Efficacy vs. answer probability shift



Step-level saliency heatmap





Chinese Reproduction: Added Qwen3-4B Runs

2 Reproduction and Extension Experiments

Same FUR setting

Dataset	Chinese C-Eval subset
Unlearning	NPO-KL, stepwise
LR / seed	10^{-5} / 1001
Specificity split	20 held-out samples
Updated params	FF2 only
CoT style	Visible Chinese CoT, no thinking mode

Runs compared


Run	CoTs	Targets	Epochs
Qwen3-8B	250	230	5
Qwen3-4B	250	230	5
Qwen3-4B	220	200	10

Important. The $n = 220$ Qwen3-4B run is not a direct duplicate: it uses **10 unlearning epochs**. We treat it as an intervention-strength extension, not as a pure model-size comparison.

Three-Run Results: Matched Five-Epoch View

2 Reproduction and Extension Experiments

Matched intervention budget: epochs 1--5



Paper-style metrics through epoch 5

Run	Eff	Spec	FF-HARD
8B, $n = 250$	63.12	96.78	7.92%
4B, $n = 250$	74.81	98.54	8.47%
4B, $n = 220$, E5 cut	79.70	97.94	11.69%


Under the matched $n = 250$, $E = 5$ setting, Qwen3-4B has much higher efficacy than Qwen3-8B, but only a small FF-HARD gain.

On the common eligible subset (158 questions), FF-HARD is 7.59% for 4B vs. 6.96% for 8B. The strongest signal is **intervention strength**, not a large faithfulness gap.

Extended 4B Run: Epoch Effect and Probability Shifts


2 Reproduction and Extension Experiments

Unlearning trajectories



E5 $\rightarrow$ E10 for 4B $n = 220$: Eff 88.99 $\rightarrow$ 90.25,
FF-HARD(all) 11.00% $\rightarrow$ 12.00%.

Endpoint FF-SOFT distribution



Ten epochs move more probability mass away from the original answer, but the gain saturates after roughly epoch 6.

Takeaway. Qwen3-4B $n = 220, E = 10$ is an **epoch-strength extension**. It supports stronger intervention effects, but should not be merged into the standard $E = 5$ model ranking.



Subsection 3

Extension Experiments1




Extension Experiments Use LoRA as a Lightweight Intervention Mechanism

2 Reproduction and Extension Experiments

- **Goal:** keep the original faithfulness-by-unlearning protocol, but replace heavier updates with LoRA adapters.
- **Backbone setting:** Phi-3 + ARC-Challenge + stepwise NPO+KL.
- **Main comparison:**
 - down_proj only
 - all-linear
 - different rank / learning-rate choices
- **Question:** can LoRA preserve strong intervention behavior while keeping specificity high?
- **Takeaway preview:** adapter target scope turns out to be the key control knob.

Question-level Metrics Show That All-linear r32 Is the Strongest Extension



2 Reproduction and Extension Experiments



- all-linear r32 reaches **44.3** Q-final FF-HARD.
- It also keeps **97.4** efficacy and **99.1** specificity.
- down_proj remains safest, but its intervention strength is much weaker.

Visual Evidence Shows a Controlled Strength–Preservation Trade-off

2 Reproduction and Extension Experiments



- all-linear r32 moves the frontier toward stronger flips while keeping specificity near 99%.
- Representative cases show clean probability redistribution rather than arbitrary collapse.

Supplementary Explorations Beyond the Main Comparison

2 Reproduction and Extension Experiments

Group	Q-FF	Eff	Spec
baseline_full	44.3	97.1	99.1
repr_only_full	44.3	97.2	99.1
firstk4_full	40.0	91.5	99.0
firstk4_repr	40.0	91.6	99.0
causal_cot_full	40.0	96.7	98.4

- **repr-loss** does not improve over the all-linear baseline.
- **first-k** still flips many questions, but weaker than full-step intervention.
- **causal-cot** changes internal behavior, but does not outperform the baseline on final flips.

Critical-step and Mechanistic Support

2 Reproduction and Extension Experiments

Group	Flip	Eff	Spec
critical baseline	22.0	98.4	99.3
critical repr	19.5	98.4	99.4
critical causal-cot	14.3	96.8	98.9

- Critical-step runs confirm that **not all reasoning steps are equally important**.
- The strongest interventions are concentrated on behaviorally important steps.
- Repr-loss and causal variants do not beat the baseline on final flip rate.

Mechanistic signal

- Causal-CoT gives the desired negative internal attention shifts.
- But better mechanistic movement does **not** automatically mean better answer-flip behavior.



Subsection 4

Extension Experiments2



Enhanced FUR: Experimental Setup

2 Reproduction and Extension Experiments

Goal

Full FUR unlearns every CoT step, which is expensive. Enhanced FUR asks whether we can unlearn only a selected top-1 step and still recover many faithful cases found by Full FUR.

Compared methods

- **Full FUR**: unlearn all segmented CoT steps.
- **Last-step top1**: unlearn only the final step.
- **Random top1**: unlearn one random step.

Run configuration

Model	Llama-3.2-3B-Instruct
Dataset	OpenBookQA
CoT cache	100 generated CoTs
Unlearning targets	80 instances
Specificity split	20 held-out samples
Method	NPO-KL, stepwise, sentencize
Epochs / LR	2 / 3e-05
Parallelism	3 shards on A100 80GB

Note. Llama-8B weights were gated at download time, so the complete run uses the locally available Llama-3B model.


Overall Results: Full FUR vs. Selective Top-1

2 Reproduction and Extension Experiments

Summary metrics

Method	Rows	Faithful inst.	Specificity
Full FUR	431	53/80 = 66.25%	94.51%
Last top1	78	22/78 = 28.21%	94.74%
Random top1	71	21/71 = 29.58%	93.73%

Full FUR finds many faithful instances, while top-1 variants use far fewer unlearning runs. All methods show high efficacy, around 89–90%, and high specificity, around 94%.




Selective FUR: Recovery–Cost Tradeoff

2 Reproduction and Extension Experiments

Recovery against Full FUR

Metric	Last top1	Random top1
Cost reduction	81.90%	83.53%
Faithful cases found	22/53	21/53
Recovery@1	41.51%	39.62%
Step-Hit@1	17.21%	18.85%
Selected-step precision	26.92%	32.39%

Top-1 selective FUR is much cheaper, but the naive selectors miss many faithful steps. Last-step is not clearly better than random, suggesting that a real verifier/ranker is needed.




Strong Positive Cases: Probability Transfer


2 Reproduction and Extension Experiments

Largest Full-FUR answer shifts

ID	Step	Change	Drop	Efficacy
508	1	$D \rightarrow A$	0.927	91.56%
266	3	$D \rightarrow C$	0.926	96.09%
1189	1	$B \rightarrow A$	0.925	92.04%
1577	7	$B \rightarrow D$	0.900	95.81%
1577	11	$B \rightarrow A$	0.894	97.16%



ID 508, step 1




ID 266, step 3

These examples show the clearest causal effect: after one reasoning step is unlearned, the original answer probability almost vanishes and another option becomes dominant.


Global Pattern and Final Takeaways

2 Reproduction and Extension Experiments

Efficacy vs. answer probability shift



Full-FUR step salience



- High efficacy does not guarantee an answer flip; probability-shift metrics are necessary.
- Full FUR is the best upper bound, but it is expensive: 431 step-runs for 80 instances.
- Top-1 selective FUR saves more than 80% of the cost, but recovers only about 40% of faithful instances.
- Next step: replace naive top-1 selection with deletion-based or verifier-guided ranking.



Outline

3 Further Exploration

- ▶ Introduction
- ▶ Reproduction and Extension Experiments
 - Reproduction
 - Chinese Reproduction
 - Extension Experiments1
 - Extension Experiments2
- ▶ Further Exploration
 - Measuring Tree of Thought Faithfulness
 - Distilled CoT via deepseek-v4-pro
- ▶ Conclusion, Limitations and Discussion



Subsection 1

Measuring Tree of Thought Faithfulness


Measuring Tree of Thought Faithfulness

3 Further Exploration

Chain-of-Thought Prompting Elicits Reasoning in Large Language Models[1]



Tree of Thoughts: Deliberate Problem Solving with Large Language Models[4]





Parametric Faithfulness: CoT vs. ToT

3 Further Exploration

Question

Does multi-path Tree-of-Thought reasoning expose a stronger dependence on the model's internal parameters than a single Chain-of-Thought path?

Core idea

- Generate reasoning traces for multiple-choice science QA.
- Unlearn selected reasoning steps while preserving held-out behavior.
- Re-evaluate whether the original answer flips after the intervention.

Runs completed

Dataset	Forget	Retain
OpenBookQA	50	20
ARC-Challenge	100	20

Model and update

- Llama-3.2-3B-Instruct
- NPO-KL unlearning
- FF2-only updates, 5 epochs
- same hyperparameters for CoT and ToT



Experimental Setup

3 Further Exploration

Reasoning arms

Arm		Construction
CoT		Single visible reasoning path.
ToT-selected		Five sampled paths; choose the highest-confidence path.
ToT	validation	sample-select vs. beam-prune before full runs.

Metrics

- **Eff**: removal of the unlearned reasoning behavior.
- **Spec**: preservation on retained examples.
- **FF-HARD**: answer flips after unlearning among eligible questions.
- **FF-SOFT**: probability shift away from the original answer.
- **Post-CoT Agree**: post-reasoning answer agreement filter.

Main Results

3 Further Exploration

Dataset	Arm	Eff	Spec	FF-HARD
OpenBookQA	CoT	81.2	95.4	72.5
OpenBookQA	ToT-selected	81.8	95.1	71.8
ARC-Challenge	CoT	75.8	95.6	68.8
ARC-Challenge	ToT-selected	79.6	95.4	72.2

Dataset	Δ Eff	Δ FF-HARD	Δ FF-SOFT
OpenBookQA	+0.6	-0.7	+5.2
ARC-Challenge	+3.8	+3.4	+7.3


Takeaways

- ARC-Challenge shows the clearest ToT advantage: higher Eff, FF-HARD, and FF-SOFT with almost unchanged Spec.
- OpenBookQA improves in soft probability shift, but not in strict paper-protocol FF-HARD.
- Paired eligible subset: ARC gains +3.2 FF-HARD points; OpenBookQA is tied.


Visual Summary

3 Further Exploration

Faithfulness comparison



Unlearning trajectories



Bottom line: ToT-selected is most useful on ARC-Challenge, where multi-path reasoning exposes stronger parameter dependence while preserving retained behavior.



Subsection 2

Distilled CoT via deepseek-v4-pro



Goal: Test Whether Distilled CoT Improves Faithfulness

3 Further Exploration

Research question

Can one SFT pass on **DeepSeek-revised CoTs** make a 3B model's generated reasoning more aligned with the parametric knowledge supporting its answers?

Base idea

- Student: LLaMA-3.2-3B-Instruct.
- Teacher revises student CoTs, not just final answers.
- Evaluate with **Faithfulness by Unlearning Reasoning Steps** (FUR).

Pipeline

1. Draft

Base LLaMA-3B generates OpenBookQA rationales.

2. Revise

deepseek-v4-pro rewrites concise, answer-consistent CoTs.

3. SFT + FUR

Fine-tune once, then compare Base vs. SFT on the same FUR test set.



Experimental Setup

3 Further Exploration

Data and SFT

Base model	Llama-3.2-3B
Teacher	deepseek-v4-pro
Draft set	2000 OpenBookQA train examples
Accepted data	SFT 1852 examples (92.6%)
SFT method	LoRA, down_proj only
LoRA config	rank 32, alpha 64, dropout 0.05
Training	lr = 10^{-4} , 1 epoch, max len 512

FUR evaluation

Target split	100 OpenBookQA test questions
Retain split	20 held-out questions
Compared arms	Base vs. DeepSeek-Rev. SFT
Unlearning	NPO + KL
FUR LR / epochs	3×10^{-5} / 5
Updated params	down_proj only
Seed	1001

Primary readout. Report FF-HARD/FF-SOFT on agreement

Main Results: Accuracy, Controls, and Faithfulness

3 Further Exploration

Model	Dir.	CoT	Agree	Steps	Eff.	Spec.
Base	70.00	73.00	82.00	6.37	92.58	94.04
SFT	73.00	77.00	85.00	3.12	97.62	98.79

Model	Agree N	FF-HARD	FF-SOFT	Valid steps
Base	82	70.73	54.09	495
SFT	85	65.88	56.07	311

What changed?

- CoTs become shorter: **6.37 to 3.12** steps.
- FUR becomes cleaner: higher efficacy and specificity.
- Faithfulness is **not clearly improved**: FF-SOFT rises, but FF-HARD falls.


Paired check. On 75 common agreement questions:

- FF-HARD: -4.00 pp, CI $[-16.00, 8.00]$.
- FF-SOFT: $+1.12$ pp, CI $[-7.95, 9.95]$.


Visual Summary

3 Further Exploration

FUR control region



Protocol faithfulness



SFT moves to the high-efficacy, high-specificity region. FF-SOFT improves slightly, while FF-HARD does not.

Takeaway. One-shot distilled-CoT SFT improves compactness and FUR control quality. It does **not yet establish a statistically stable faithfulness gain.**



Outline

4 Conclusion, Limitations and Discussion

- ▶ Introduction
- ▶ Reproduction and Extension Experiments
 - Reproduction
 - Chinese Reproduction
 - Extension Experiments1
 - Extension Experiments2
- ▶ Further Exploration
 - Measuring Tree of Thought Faithfulness
 - Distilled CoT via deepseek-v4-pro
- ▶ Conclusion, Limitations and Discussion



Key Conclusions

4 Conclusion, Limitations and Discussion

- **FUR Method:** The first machine unlearning-based tool for measuring parametric faithfulness of Chain-of-Thought (CoT)
- **Core Finding:** Some CoTs contain reasoning steps that genuinely influence model decisions
- **Method Comparison:** Parameter-level intervention outperforms traditional context perturbation methods
- **Mechanistic Validation:** Unlearning alters internal reasoning paths rather than merely perturbing answers
- **Key Implication:** Faithfulness and human interpretability require separate optimization



Limitations

4 Conclusion, Limitations and Discussion

- **Scope Restriction:** Only analyzes samples with consistent CoT/no-CoT answers, reducing contextual interference but limiting applicability
- **Unlearning Imperfection:** Relies on machine unlearning, which may produce incomplete or imprecise results that compromise faithfulness assessments
- **Closed-Source Incompatibility:** Requires model parameter access and fine-tuning, cannot be directly applied to API-only models
- **Task Limitation:** Experiments are limited to English MCQA tasks; extension to long-form generation remains unvalidated
- **High Computational Cost:** Each CoT step needs separate unlearning, scaling poorly with long reasoning chains (e.g., o1/R1)



Discussion

4 Conclusion, Limitations and Discussion

- **Stability:** How can we adjust the learning rate and other hyperparameters to achieve more stable results?
- **Explainability:** The explainability of the final results is determined by the explainability of large language models and Chain-of-Thought.
- **Improvement :** How to improve the parametric faithfulness of CoT



References

4 Conclusion, Limitations and Discussion

- [1] Jason Wei et al. “Chain of Thought Prompting Elicits Reasoning in Large Language Models”. In: (). arXiv: 2201.11903.
- [2] Tamera Lanham et al. Measuring Faithfulness in Chain-of-Thought Reasoning. arXiv: 2307.13702 [cs.AI].
- [3] Martin Tutek et al. “Measuring Chain of Thought Faithfulness by Unlearning Reasoning Steps”. In: ed. by Christos Christodoulopoulos et al. Suzhou, China: Association for Computational Linguistics.
- [4] Shunyu Yao et al. “Tree of Thoughts: Deliberate Problem Solving with Large Language Models”. In.

et al. (Due to space constraints, only the above are shown.)



Github Link

4 Conclusion, Limitations and Discussion

Measuring Faithfulness Group 27

<https://github.com/KouseiAimer/Measuring-Faithfulness-Group27>
github.com/KouseiAimer/Measuring-Faithfulness-Group27

More Faithful CoT (Original v2)

<https://github.com/tyyyyyy333/more-faithful-cot>
github.com/tyyyyyy333/more-faithful-cot



Measuring Chain of Thought Faithfulness

Thank You for Listening!